

## Homework 2 - IEEE Fraud Detection

For all parts below, answer all parts as shown in the Google document for Homework 2. Be sure to include both code that justifies your answer as well as text to answer the questions. We also ask that code be commented to make it easier to follow.

```
In [0]: #Drive mounted for reading files.
from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response\\_type=code](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aoob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code)

Enter your authorization code:

.....

Mounted at /content/drive

```
In [0]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn import preprocessing
from scipy.stats import spearmanr
from scipy.stats import pearsonr
```

```
In [0]: #Reading files with pandas
train_transaction = pd.read_csv("/content/drive/My Drive/train_transaction.csv", index_col= 'TransactionID')
train_identity = pd.read_csv("/content/drive/My Drive/train_identity.csv", index_col= 'TransactionID')
test_transaction = pd.read_csv("/content/drive/My Drive/test_transaction.csv", index_col= 'TransactionID')
test_identity = pd.read_csv("/content/drive/My Drive/test_identity.csv", index_col= 'TransactionID')

print("File Reading Successfull!!")
```

File Reading Successfull!!

In [0]: *#Checking how many transactions have associated identity information.*

```
associated_train_data = np.sum(train_transaction.index.isin(train_identity.index.unique()))
associated_test_data = np.sum(test_transaction.index.isin(test_identity.index.unique()))

#Percentage of associated transaction-identity information.
train_association = associated_train_data/len(train_transaction.index)*100
test_association = associated_test_data/len(test_transaction.index)*100

print(train_association,'%','have associated identity information in training data')
print(test_association,'%','have associated identity information in testing data')
```

24.42391709283029 % have associated identity information in training data  
28.006615471756945 % have associated identity information in testing data

In [0]: *#Merging transaction*  
*#Merging with how = "left", guarantees that there are still 590540 training records and 506691 testing records,*  
*#even though not every transaction has associated identity information.*  
train\_data = train\_transaction.merge(train\_identity, how = "left", left\_index = True, right\_index = True)  
test\_data = test\_transaction.merge(test\_identity, how = "left", left\_index = True, right\_index = True)  
  
print(train\_data.info())  
print(test\_data.info())

*#For the next few cells, I will be cleaning and transforming the data.*

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 590540 entries, 2987000 to 3577539
Columns: 433 entries, isFraud to DeviceInfo
dtypes: float64(385), int64(17), object(31)
memory usage: 1.9+ GB
None
<class 'pandas.core.frame.DataFrame'>
Int64Index: 506691 entries, 3663549 to 4170239
Columns: 432 entries, TransactionDT to DeviceInfo
dtypes: float64(399), int64(2), object(31)
memory usage: 1.7+ GB
None
```

```
In [0]: #Before we get rid of the NaN values, let's first find the percentage of missing values in each column.
train_data_missing = train_data.isna()
percentage_miss = train_data_missing.sum()/len(train_data)*100

print(percentage_miss)
#Now we identify which columns have more than 50% of data missing. i.e. NaN values
percentage_miss_filter = percentage_miss > 50
#print(percentage_miss_filter)

#We will use this info later to filter the features for our prediction model.
```

|                |           |
|----------------|-----------|
| isFraud        | 0.000000  |
| TransactionDT  | 0.000000  |
| TransactionAmt | 0.000000  |
| ProductCD      | 0.000000  |
| card1          | 0.000000  |
| card2          | 1.512683  |
| card3          | 0.265012  |
| card4          | 0.267044  |
| card5          | 0.721204  |
| card6          | 0.266028  |
| addr1          | 11.126427 |
| addr2          | 11.126427 |
| dist1          | 59.652352 |
| dist2          | 93.628374 |
| P_emaildomain  | 15.994852 |
| R_emaildomain  | 76.751617 |
| C1             | 0.000000  |
| C2             | 0.000000  |
| C3             | 0.000000  |
| C4             | 0.000000  |
| C5             | 0.000000  |
| C6             | 0.000000  |
| C7             | 0.000000  |
| C8             | 0.000000  |
| C9             | 0.000000  |
| C10            | 0.000000  |
| C11            | 0.000000  |
| C12            | 0.000000  |
| C13            | 0.000000  |
| C14            | 0.000000  |
| ...            |           |
| id_11          | 76.127273 |
| id_12          | 75.576083 |
| id_13          | 78.440072 |
| id_14          | 86.445626 |
| id_15          | 76.126088 |
| id_16          | 78.098012 |
| id_17          | 76.399736 |
| id_18          | 92.360721 |
| id_19          | 76.408372 |
| id_20          | 76.418024 |
| id_21          | 99.126393 |
| id_22          | 99.124699 |
| id_23          | 99.124699 |
| id_24          | 99.196159 |
| id_25          | 99.130965 |
| id_26          | 99.125715 |
| id_27          | 99.124699 |
| id_28          | 76.127273 |
| id_29          | 76.127273 |
| id_30          | 86.865411 |
| id_31          | 76.245132 |
| id_32          | 86.861855 |
| id_33          | 87.589494 |
| id_34          | 86.824771 |
| id_35          | 76.126088 |
| id_36          | 76.126088 |

```
id_37      76.126088
id_38      76.126088
DeviceType  76.155722
DeviceInfo  79.905510
Length: 433, dtype: float64
```

```
In [0]: #Replacing NaN values with the most frequent value in the column.
train_data = train_data.fillna(train_data.mode().iloc[0])
test_data = test_data.fillna(test_data.mode().iloc[0])
print(train_data.head())
print(test_data.head())
```

```
      isFraud  ...      DeviceInfo
TransactionID  ...
2987000      0  ...      Windows
2987001      0  ...      Windows
2987002      0  ...      Windows
2987003      0  ...      Windows
2987004      0  ...  SAMSUNG SM-G892A Build/NRD90M
```

[5 rows x 433 columns]

```
      TransactionDT  TransactionAmt  ... DeviceType  DeviceInfo
TransactionID      ...
3663549      18403224      31.95  ...  desktop      Windows
3663550      18403263      49.00  ...  desktop      Windows
3663551      18403310      171.00  ...  desktop      Windows
3663552      18403310      284.95  ...  desktop      Windows
3663553      18403317      67.95  ...  desktop      Windows
```

[5 rows x 432 columns]

```
In [0]: #These are the features we will consider in our prediction model
train_data_final = train_data.filter(['TransactionID', 'TransactionDT', 'TransactionAmt', 'ProductCD', 'DeviceType', 'DeviceInfo',
                                     'addr1', 'addr2', 'card1', 'card2', 'card3', 'card4', 'card5', 'card6', 'P_emaildomain', 'C1', 'C2', 'C3', 'C4', 'C5', 'C6',
                                     'C7', 'C8', 'C9', 'C10', 'C11', 'C12', 'C13', 'C14', 'isFraud'])

test_data_final = test_data.filter(['TransactionID', 'TransactionDT', 'TransactionAmt', 'ProductCD', 'DeviceType', 'DeviceInfo',
                                   'addr1', 'addr2', 'card1', 'card2', 'card3', 'card4', 'card5', 'card6', 'P_emaildomain', 'C1', 'C2', 'C3', 'C4', 'C5', 'C6',
                                   'C7', 'C8', 'C9', 'C10', 'C11', 'C12', 'C13', 'C14'])
```

```
In [0]: #Separating the features and the labels
features_train = train_data_final.drop('isFraud', axis = 1)
label_train = train_data_final['isFraud']
features_test = test_data_final.copy()
```

```
In [0]: #Mapping categorical data to numeric values
for feature in features_train:
    if features_train[feature].dtype == 'object' or features_test[feature].dtype == 'object':
        le = preprocessing.LabelEncoder()
        le.fit(list(features_train[feature].values) + list(features_test[feature].values))
        features_train[feature] = le.transform(list(features_train[feature].values))
        features_test[feature] = le.transform(list(features_test[feature].values))
```

```
In [0]: #Filtering data for parts 1 to 5.

train_data_filtered = train_data.filter(['TransactionID', 'DeviceType', 'DeviceInfo', 'TransactionDT', 'TransactionAmt', 'ProductCD',
                                         'addr1', 'addr2', 'card4', 'card6', 'dist1', 'dist2', 'P_emaildomain', 'R_emaildomain', 'isFraud'])
```


```
In [0]: train_data_filtered.head()
```

## Part 1 - Fraudulent vs Non-Fraudulent Transaction

```
In [0]: # TODO: code and runtime results

#Separating the fraudulent and Non-fraudulent transactions
train_transaction_fraud = train_data_filtered[train_data_filtered['isFraud'] == 1]
train_transaction_notFraud = train_data_filtered[train_data_filtered['isFraud'] == 0]
```

```
In [79]: #
figure = plt.figure(figsize=(5,5))
plt.xlim(70,78)
plt.ylim(0,1000)
plt.hist(train_transaction_fraud['TransactionDT']/86400, bins = 1800, alpha =
1, label = 'Fraud', color = 'red')
plt.hist(train_transaction_notFraud['TransactionDT']/86400, bins = 1800, alpha
= 0.2, label = 'Not Fraud', color = 'blue')
plt.legend(loc='upper right')
plt.xlabel("Days")
plt.ylabel("No of Transactions")
plt.show()
```




### Answer:

In the above graph, I have plotted the number of fraudulent (red) and non-fraudulent (purple) transactions against day of the week.

**There is no significant correlation between the day of the week and whether a transaction is fraudulent or not.**

```
In [0]: #
plt.hist(train_transaction_fraud['TransactionAmt'], alpha = 1, label = 'Fraud',
, color = 'red')
plt.hist(train_transaction_notFraud['TransactionAmt'], alpha = 0.4, label = 'N
ot Fraud', color = 'yellow')
plt.legend(loc='upper right')
plt.xlabel("Amount")
plt.ylabel("No of Transactions")
plt.xlim(0,3500,500)
plt.show()
```

**Answer:**

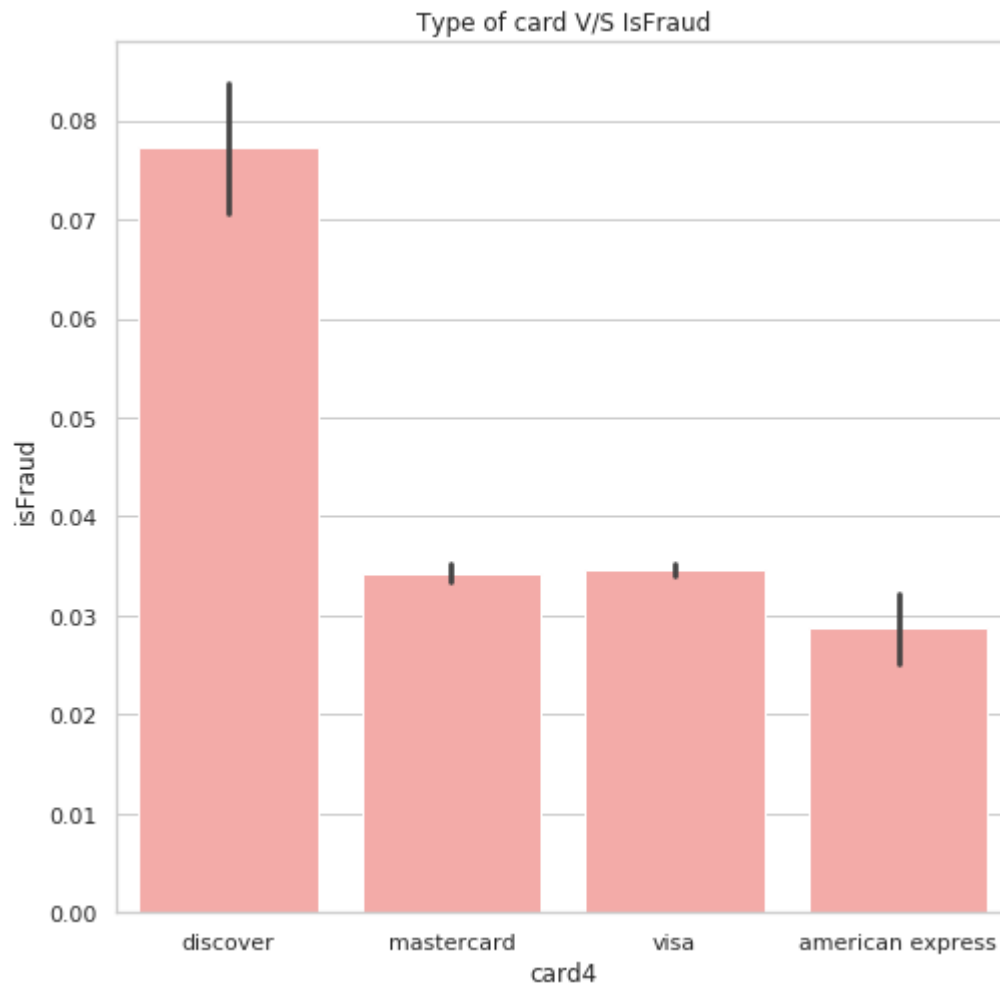
In the above graph, I have plotted the Transaction amounts, number of transactions and whether they are fraudulent or not.

**As we can see, the fraudulent (red) transactions are mostly of lower amounts between 0-500.**




```
In [82]: #Checking which type of card saw highest number of fraudulent transactions.
sns.set(style="whitegrid")
figure = plt.figure(figsize=(8,8))
sns.set_color_codes("pastel")
sns.barplot(x=train_data_filtered['card4'],y=train_data_filtered['isFraud'],
            data=train_data_filtered,label="Fraud or not",color = "r").set_title("Type of card V/S IsFraud")
```

Out[82]: Text(0.5, 1.0, 'Type of card V/S IsFraud')




```
In [66]: #Checking which type of card saw highest number of fraudulent transactions.
sns.set(style="whitegrid")
figure = plt.figure(figsize=(8,8))
sns.set_color_codes("dark")
sns.barplot(x=train_data_filtered['card6'],y=train_data_filtered['isFraud'],
            data=train_data_filtered,label="Fraud or not",color = "g").set_title(
            "Type of transaction V/S IsFraud")
```

Out[66]: Text(0.5, 1.0, 'Type of transaction V/S IsFraud')



```
In [85]: #Checking which type of operating system saw highest number of fraudulent transactions.
sns.set(style="whitegrid")
figure = plt.figure(figsize=(8,8))
sns.set_color_codes("pastel")
list = train_transaction_fraud['DeviceInfo'].value_counts().nlargest(20)
sns.barplot(x=list.values,y=list.index,data=train_transaction_fraud,label="Fraud or not",color = "g").set_title("Type of operating system V/S IsFraud")
```

Out[85]: Text(0.5, 1.0, 'Type of operating system V/S IsFraud')



### Answer:

The graph of "Type of Operating system v/s isFraud" helps us identify what number of transactions done on a particular operating systems were fraudulent.

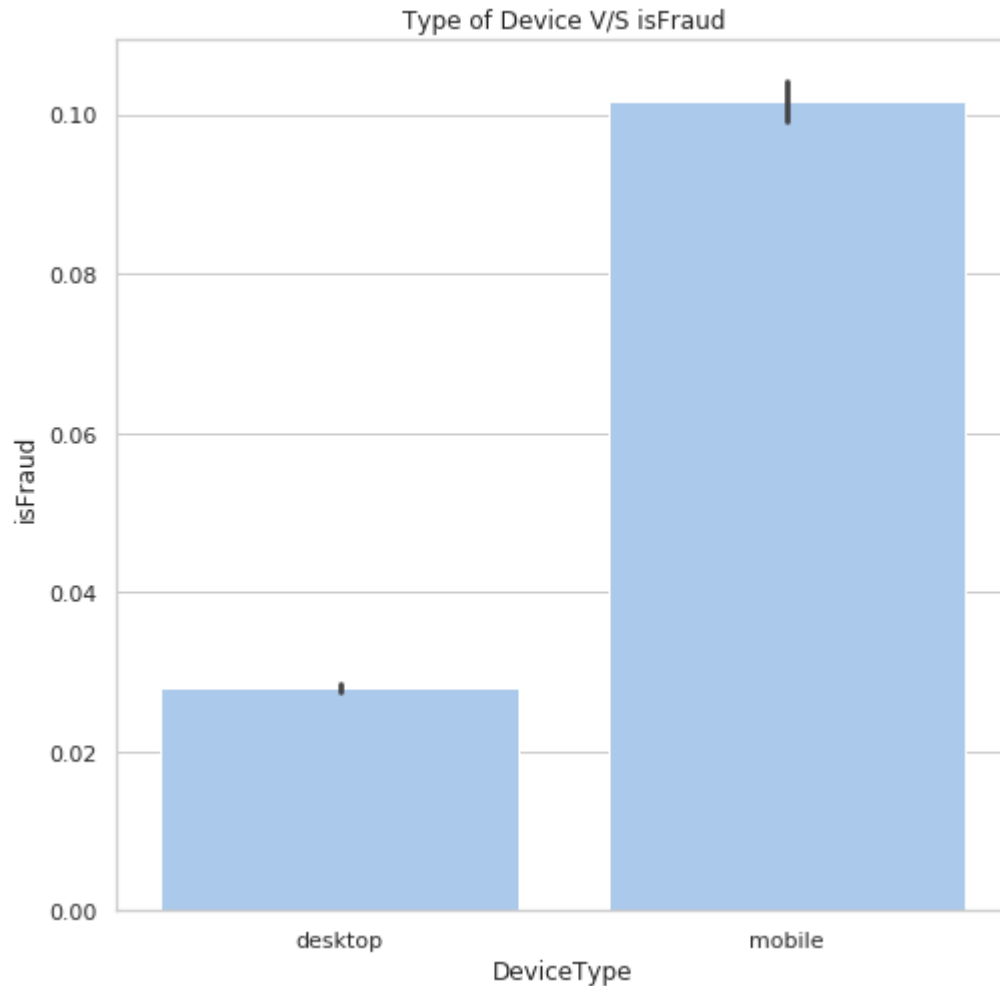
E.g.:-

Windows: >14000

MacOS: < 100

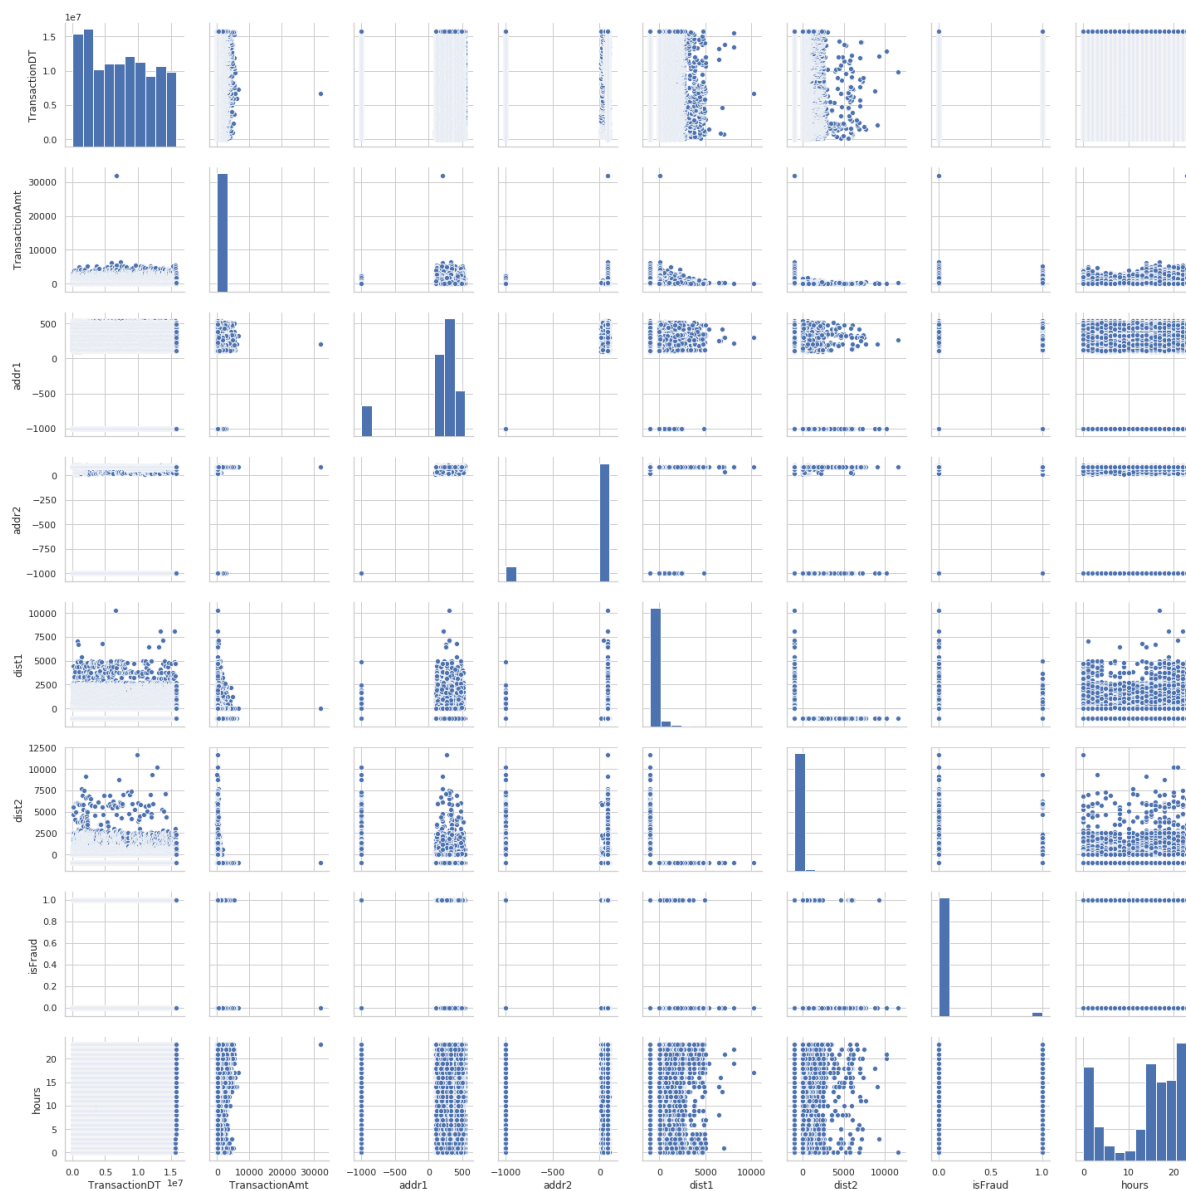
```
In [83]: #Checking which type of device saw highest number of fraudulent transactions.
sns.set(style="whitegrid")
figure = plt.figure(figsize=(8,8))
sns.set_color_codes("pastel")
sns.barplot(x=train_data_filtered['DeviceType'],y=train_data_filtered['isFraud'],
            data=train_data_filtered,label="Fraud or not",color = "b").set_title(
            "Type of Device V/S isFraud")
```

Out[83]: Text(0.5, 1.0, 'Type of Device V/S isFraud')



```
In [0]: sns.pairplot(train_data_filtered)
```

```
Out[0]: <seaborn.axisgrid.PairGrid at 0x7ff16591d588>
```



```
In [87]: fig, ax = plt.subplots(2, 2, figsize = (10,10))

time = train_transaction_fraud['TransactionDT']
time2 = train_transaction_notFraud['TransactionDT']

log_time = np.log(time.values/86400)
log_time2 = np.log(time2.values/86400)

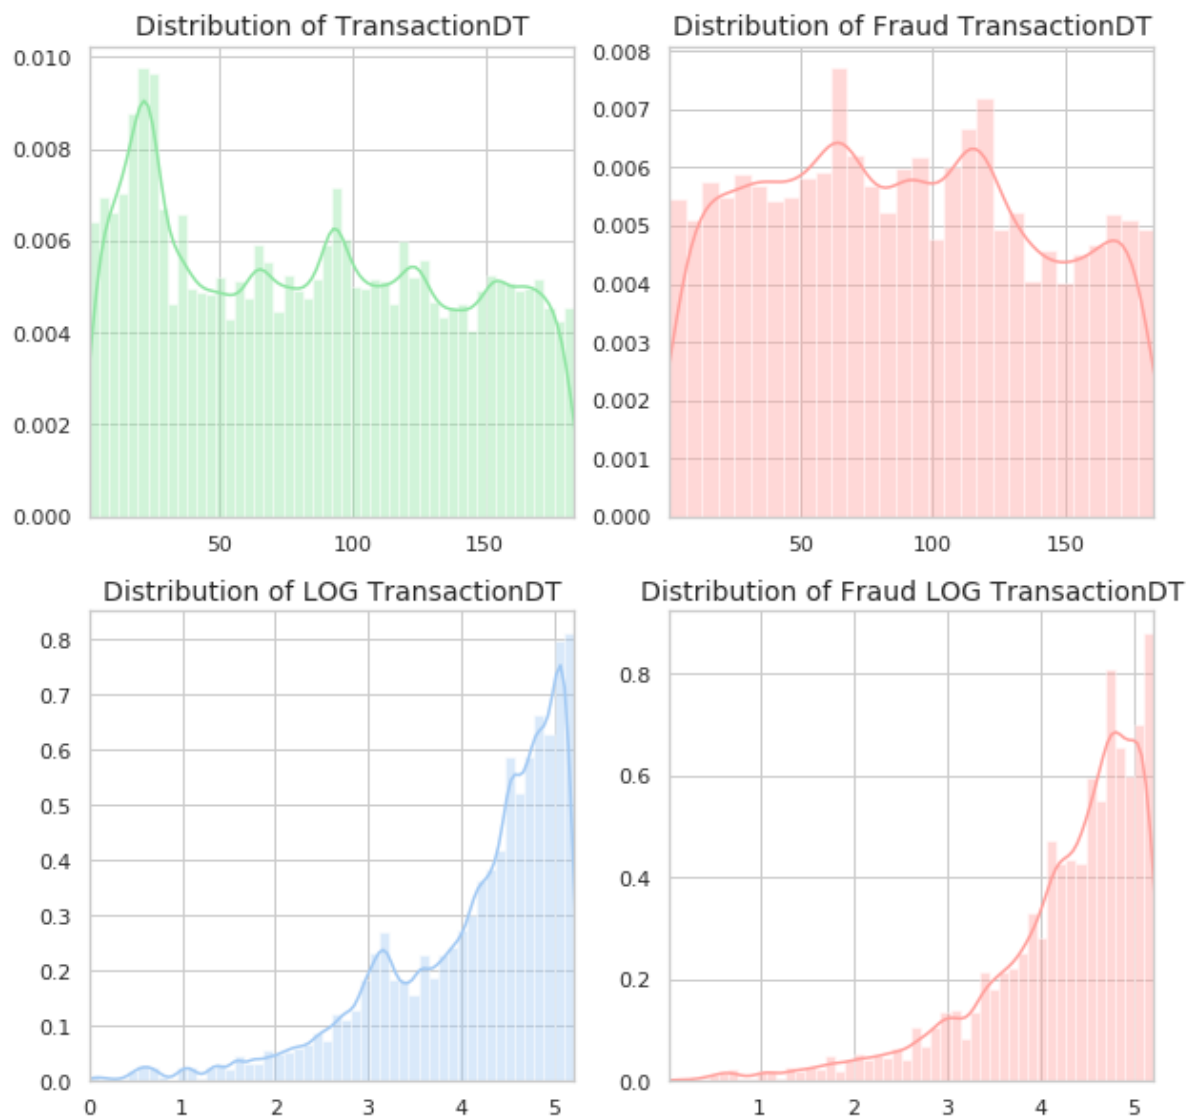
sns.distplot(time2.values/86400, ax = ax[0,0], color = 'g')
ax[0,0].set_title('Distribution of TransactionDT', fontsize=14)
ax[0,0].set_xlim([min(time2.values/86400), max(time2.values/86400)])

sns.distplot(time.values/86400, ax=ax[0,1], color='r')
ax[0,1].set_title('Distribution of Fraud TransactionDT', fontsize=14)
ax[0,1].set_xlim([min(time.values/86400), max(time.values/86400)])

sns.distplot(log_time2, ax=ax[1,0], color='b')
ax[1,0].set_title('Distribution of LOG TransactionDT', fontsize=14)
ax[1,0].set_xlim([min(log_time2), max(log_time2)])

sns.distplot(log_time, ax=ax[1,1], color='r')
ax[1,1].set_title('Distribution of Fraud LOG TransactionDT', fontsize=14)
ax[1,1].set_xlim([min(log_time), max(log_time)])

plt.show()
```



**Answer:**

In the above graph, I have plotted the transaction date time V/S Fraud or Not.

## Part 2 - Transaction Frequency

```
In [0]: # TODO: code to generate the frequency graph

def make_hour_feature(df, tname='TransactionDT'):
    """
    Creates an hour of the day feature, encoded as 0-23.

    Parameters:
    -----
    df : pd.DataFrame
        df to manipulate.
    tname : str
        Name of the time column in df.
    """
    hours = df[tname] / (3600)
    encoded_hours = np.floor(hours) % 24
    return encoded_hours


mode = np.max(train_data_filtered.addr2.mode().iloc[0])
data_addr = train_data_filtered.loc[train_data_filtered['addr2'] == mode]
data_addr['hours'] = make_hour_feature(data_addr)
data_addr['hours'].hist(bins = 24 , xrot = 30, color = "skyblue", lw = 1)
plt.title("Transaction Frequency per hour")
plt.ylabel("No. of transactions")
plt.xlabel("Hour")
```

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:19: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

Out[0]: Text(0.5, 0, 'Hour')





**Answer:**

From the 24 hours in a day (shown as 0-23 in the graph above), one can clearly see the dip in the graph of number of transactions per hour. The dip represents the sleeping hours of the vast majority. An inverse bell curve graph helps determine the waking v/s sleeping hours for the most frequent country code which is 87.0

**Part 3 - Product Code**

```
In [0]: # TODO: code to analyze prices for different product codes

maxCost = 0
minCost = 100
for i in train_data_filtered.ProductCD.unique():
    print(i)
    print("Count: ",train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.count())
    print("Mean: ",train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.mean())
    count = train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.count()
    median = train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.median()
    mean = train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.mean()
    std_deviation = train_data_filtered[train_data_filtered.ProductCD == i].TransactionAmt.std()
    if mean > maxCost:
        maxCost = mean
        product = i

print("The most expensive product is",product,"and it's average cost is",maxCost)

for j in train_data_filtered.ProductCD.unique():
    mean = train_data_filtered[train_data_filtered.ProductCD == j].TransactionAmt.mean()
    if mean < minCost:
        minCost = mean
        product = j

print("The cheapest product is",product,"and it's average cost is",minCost)

figure = plt.figure(figsize=(8,8))
#ax = sns.boxplot( x=train_data_filtered['ProductCD'], y=train_data_filtered['TransactionAmt'], hue = train_data_filtered['isFraud'], palette = "Set2")
ax = sns.boxplot( x=train_data_filtered['ProductCD'], y=np.log(train_data_filtered['TransactionAmt']), hue = train_data_filtered['isFraud'], palette = "Set2")
ax.legend(frameon=False, loc='upper right', ncol=1)
```

W

Count: 439670

Mean: 153.15855385223293

H

Count: 33024

Mean: 73.17005813953489

C

Count: 68519

Mean: 42.872353113733446

S

Count: 11628

Mean: 60.269487444100434

R


Count: 37699

Mean: 168.30618849306347

The most expensive product is R and it's average cost is 168.30618849306347

The cheapest product is C and it's average cost is 42.872353113733446

Out[0]: <matplotlib.legend.Legend at 0x7f7471a96048>



**Answer:**

The average cost of each product is calculated by the total cost of that product divided by the total number of product transactions. By calculating the above, we have found that Product **R** is the **most expensive** product with an average cost of **168.30** and Product **C** is the **cheapest** one with an average cost of **42.872**.


The graph above shows product W as the most expensive, however, that is just because it has more number of transactions and possibly more outliers. To avoid this, I have used mean average cost as the metric.

## Part 4 - Correlation Coefficient

```
In [0]: # TODO: code to calculate correlation coefficient
from datetime import datetime
train_data_filtered['TransactionTime'] = pd.to_datetime(train_data_filtered['T
ransactionDT'],unit = 's').dt.time


#Plot to visualize date time v/s number of transactions
figure = plt.figure(figsize=(8,4))
vals = plt.hist(train_data_filtered['TransactionDT'] / (3600*24), bins=1800)
plt.xlim(70, 78)
plt.xlabel('Days')
plt.ylabel('Number of transactions')
plt.ylim(0,1000)
```

Out[0]: (0, 1000)



```
In [0]: #By grouping them in hours and then summing the amount of transactions per hour gives below
train_data_filtered['hours'] = make_hour_feature(train_data_filtered)
dataGroupedByHours = train_data_filtered.groupby('hours')
plt.scatter(dataGroupedByHours['hours'].unique(),dataGroupedByHours.TransactionAmt.sum())
```

Out[0]: <matplotlib.collections.PathCollection at 0x7f746f497f28>



```
In [0]: print("Spearman Correlation coefficient with sum of transaction amounts: ",spearmanr(dataGroupedByHours['hours'].unique(),dataGroupedByHours.TransactionAmt.sum()))
```

Spearman Correlation coefficient with sum of transaction amounts: SpearmanrResult(correlation=0.6304347826086956, pvalue=0.0009590059599017164)


### Answer:

The correlation coefficient between hours of a day and the sum of transaction amount per hour is 0.63. This is a pretty strong correlation which signifies how hours (waking v/s sleeping) can affect the frequency of large transactions made.

## Part 5 - Interesting Plot

```
In [0]: plt.figure(figsize=(10,8))
sns.heatmap(train_data_final.corr())
```

```
Out[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd87bb97898>
```




### Answer:

The above heatmap helps us understand the correlation between the various columns of our dataset. However, a quick glance tells us that even though most of them depict very low correlation with each other, **the attribute 'card3' is highly correlated with most of the other attributes and it shows a correlation of around 0.2 with our target attribute 'isFraud'**. Some interesting insights can be found with this information and our model predicts better when this column is included.

```
In [88]: #No overlap between train and test data dates.
train_data['TransactionDT'].plot(kind = "hist", label = "Train", bins = 50)
test_data['TransactionDT'].plot(kind = "hist", label = "Test", bins = 50)
plt.legend()
```

Out[88]: <matplotlib.legend.Legend at 0x7fd8755e8fd0>



### Answer:

The above graph shows that there is no overlap between the training and testing data. This signifies that the transactions recorded for training and testing were not done during the same time interval but rather some time intervals apart.

**Train:** min = 86400, max = 15811131

**Test:** min = 18403224, max = 34214345

If we assume TransactionDT is in seconds, then:

Time span of the total dataset is 394.9993634259259 days.

Time span of Train dataset is 181.99920138888888 days.

Time span of Test dataset is 182.99908564814814 days.

The gap between train and test is 30.00107638888889 days.

## Part 6 - Prediction Model

```
In [0]: #Definig parameters for my model.
x_train = features_train
y_train = label_train
x_test = features_test
from sklearn.model_selection import train_test_split
data_train, data_test, target_train, target_test = train_test_split(x_train,y_
train, test_size = 0.30, random_state = 10)
```



```
In [77]: # TODO: code for your final model

#Random Forests Classifier
from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier()
clf.fit(x_train, y_train)
p_prob = clf.predict_proba(x_test)
preds = clf.predict(x_test)
sub = pd.read_csv('/content/drive/My Drive/sample_submission.csv', index_col=
'TransactionID')
sub['isFraud'] = preds
sub.to_csv('myPrediction17.csv')

#Measuring the accuracy of my model.
from sklearn.metrics import accuracy_score
#Using the object of Random Forests classifier
#train the algorithm on training data and predict using the testing data
pred = clf.fit(data_train, target_train).predict(data_test)
#print the accuracy score of the model
accuracy_score = accuracy_score(target_test, pred, normalize = True)


print("Random Forests - Percentage accuracy: ",accuracy_score * 100)

#Plotting the probabilities
pd.DataFrame(p_prob).hist()
plt.xlabel("Probabilites", ha = 'center')
plt.ylabel("Number of transactions", va = 'center')
```

/usr/local/lib/python3.6/dist-packages/sklearn/ensemble/forest.py:245: Future Warning: The default value of n\_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

Out[77]: Text(0, 0.5, 'Number of transactions')



```
In [78]: #Model2
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
logreg.fit(x_train, y_train)
p_prob = logreg.predict_proba(x_test)
y_pred = logreg.predict(x_test)
sub = pd.read_csv('/content/drive/My Drive/sample_submission.csv', index_col=
'TransactionID')
sub['isFraud'] = y_pred
sub.to_csv('myPrediction18.csv')

#Measuring the accuracy of my model.
from sklearn.metrics import accuracy_score
#Using the object of Random Forests classifier
#train the algorithm on training data and predict using the testing data
pred2 = logreg.fit(data_train, target_train).predict(data_test)
#print the accuracy score of the model
accuracy_score = accuracy_score(target_test, pred2, normalize = True)

print("Logistic Regression - Percentage accuracy: ",accuracy_score * 100)

#Plotting the probabilities
pd.DataFrame(p_prob).hist()
plt.xlabel("Probabilites", ha = 'center')
plt.ylabel("Number of transactions", va = 'center')
```


```

/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a s
olver to silence this warning.
FutureWarning)
/usr/local/lib/python3.6/dist-packages/sklearn/linear_model/logistic.py:432:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a s
olver to silence this warning.
FutureWarning)

```

Logistic Regression - Percentage accuracy: 96.44393267179193

Out[78]: Text(0, 0.5, 'Number of transactions')



### Answer:

I have worked on two models for predicting whether a transaction is fraudulent or not - Random forests classifier and a logistic regression model. The random forests one works better as it uses a modified tree learning algorithm, where in at each step of the learning process, it selects a random subset of features. This way it tries out different combinations of the feature sets and finally ensembles the ones which give out the most accurate predictions.

**Accuracy of this model:** 98.45%

**Kaggle score:** 0.7089

## Part 7 - Final Result

Report the rank, score, number of entries, for your highest rank. Include a snapshot of your best score on the leaderboard as confirmation. Be sure to provide a link to your Kaggle profile. Make sure to include a screenshot of your ranking. Make sure your profile includes your face and affiliation with SBU.

Kaggle Link: [PayalMehta\\_Kaggle\\_Link \(https://www.kaggle.com/payal95/competitions\)](https://www.kaggle.com/payal95/competitions)

Highest Rank: 5532

Score: 0.7089

Number of entries: 19

#### Credits and References:

1. [Time and Day - Predictive feature \(https://www.kaggle.com/fchmiel/day-and-time-powerful-predictive-feature\)](https://www.kaggle.com/fchmiel/day-and-time-powerful-predictive-feature)
2. [Fraud\\_models \(https://www.kaggle.com/jesucristo/fraud-complete-eda/notebook#Models\)](https://www.kaggle.com/jesucristo/fraud-complete-eda/notebook#Models)

```
In [71]: #@title Part 8 - Kaggle Rank
%%html
<iframe src = "https://drive.google.com/uc?id=10GNtjUoniouTm7r6nmVaeI1Nv4RGzCB
d" width = "840" height = "480"></iframe>
```

| Overview          | Data            | Notebooks | Discussion | Leaderb |
|-------------------|-----------------|-----------|------------|---------|
| 5529              | Volkmar         |           |            |         |
| 5530              | shuhum          |           |            |         |
| 5531              | Steve RASSINOT  |           |            |         |
| 5532              | Payal Mehta     |           |            |         |
| Your Best Entry ↑ |                 |           |            |         |
| 5533              | DrPurshottam KH |           |            |         |